

Student name: Roja Kamble, 11454258

Assignment: 4

Part 1: Data Wrangling

You have to write code to answer the questions below

- Import pandas library
- Read the data stored in your local machine <https://www.kaggle.com/datasets/yasserh/student-marks-dataset>
(<https://www.kaggle.com/datasets/yasserh/student-marks-dataset>)
- Name the variable df

Show information about the dataset

```
In [6]: import pandas as pd
import numpy as np

df = pd.read_csv( '/Users/roja/Downloads/Student_Marks.csv' )
df
```

Out[6]:

	number_courses	time_study	Marks
0	3	4.508	19.202
1	4	0.096	7.734
2	4	3.133	13.811
3	6	7.909	53.018
4	8	7.811	55.299
...
95	6	3.561	19.128
96	3	0.301	5.609
97	4	7.163	41.444
98	7	0.309	12.027
99	3	6.335	32.357

100 rows × 3 columns

Show the last 10 rows

```
In [7]: rows = df.tail(10)
rows
```

Out[7]:

	number_courses	time_study	Marks
90	7	4.182	24.394
91	8	2.730	19.564
92	4	5.027	23.916
93	8	6.471	42.426
94	8	3.919	24.451
95	6	3.561	19.128
96	3	0.301	5.609
97	4	7.163	41.444
98	7	0.309	12.027
99	3	6.335	32.357

Reset the index such that it starts from 1 (instead of 0)

```
In [8]: df.index = np.arange(1, len(df) + 1)
df
```

Out[8]:

	number_courses	time_study	Marks
1	3	4.508	19.202
2	4	0.096	7.734
3	4	3.133	13.811
4	6	7.909	53.018
5	8	7.811	55.299
...
96	6	3.561	19.128
97	3	0.301	5.609
98	4	7.163	41.444
99	7	0.309	12.027
100	3	6.335	32.357

100 rows × 3 columns

Lowercase all columns

```
In [9]: # Basically column heads are already lowercase except 'Marks' column
```

```
t1 = pd.read_csv('/Users/roja/Downloads/Student_Marks.csv')
t1.head()
t1.columns = t1.columns.str.lower()
t1
```

```
Out[9]:
```

	number_courses	time_study	marks
0	3	4.508	19.202
1	4	0.096	7.734
2	4	3.133	13.811
3	6	7.909	53.018
4	8	7.811	55.299
...
95	6	3.561	19.128
96	3	0.301	5.609
97	4	7.163	41.444
98	7	0.309	12.027
99	3	6.335	32.357

100 rows × 3 columns

In [10]: *# Uppercase all columns*

```
t2 = pd.read_csv('/Users/roja/Downloads/Student_Marks.csv')
t2.head()
t2.columns = t2.columns.str.upper()
t2
```

Out[10]:

	NUMBER_COURSES	TIME_STUDY	MARKS
0	3	4.508	19.202
1	4	0.096	7.734
2	4	3.133	13.811
3	6	7.909	53.018
4	8	7.811	55.299
...
95	6	3.561	19.128
96	3	0.301	5.609
97	4	7.163	41.444
98	7	0.309	12.027
99	3	6.335	32.357

100 rows × 3 columns

Return a boolean value indicating whether the dataset has missing values

```
In [11]: df.isnull()
```

```
Out[11]:
```

	number_courses	time_study	Marks
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
5	False	False	False
...
96	False	False	False
97	False	False	False
98	False	False	False
99	False	False	False
100	False	False	False

100 rows × 3 columns

Show average of all columns grouped by number_course in a same DataFrame table

```
In [12]: avg = df.groupby([ 'number_courses' ]).mean()  
avg
```

Out[12]:

	time_study	Marks
number_courses		
3	3.685773	18.433318
4	3.598524	19.029952
5	3.070600	17.641000
6	4.838750	29.863062
7	4.242667	29.401467
8	4.955750	33.835375

Show the sum of time_study and mean of marks grouped by each number_course in a same DataFrame table

```
In [13]: df.mean()
```

```
Out[13]: number_courses    5.29000  
time_study    4.07714  
Marks    24.41769  
dtype: float64
```



```
In [14]: sum = df.groupby(['number_courses'])['time_study'].sum()  
sum
```

```
Out[14]: number_courses  
3      81.087  
4      75.569  
5      30.706  
6      77.420  
7      63.640  
8      79.292  
Name: time_study, dtype: float64
```

```
In [15]: mean = df.groupby(['number_courses'])['Marks'].mean()  
mean
```

```
Out[15]: number_courses  
3      18.433318  
4      19.029952  
5      17.641000  
6      29.863062  
7      29.401467  
8      33.835375  
Name: Marks, dtype: float64
```

```
In [16]: pip install -U prettytable
```

```
Requirement already satisfied: prettytable in /Users/roja/opt/anaconda3/lib/python3.9/site-packages (3.5.0)  
Requirement already satisfied: wcwidth in /Users/roja/opt/anaconda3/lib/python3.9/site-packages (from prettytable) (0.2.5)  
Note: you may need to restart the kernel to use updated packages.
```

```
In [17]: from prettytable import PrettyTable

res = PrettyTable([sum, mean])
res
```

```
Out[17]:
```

	number_courses	number_courses
	3 81.087	3 18.433318
	4 75.569	4 19.029952
	5 30.706	5 17.641000
	6 77.420	6 29.863062
	7 63.640	7 29.401467
	8 79.292	8 33.835375

Name: time_study, dtype: float64 Name: Marks, dtype: float64

Write a class My_KMeans that implements k-means clustering algorithm. You are required to have the following attributes

- Parameters:
 - Default clusters
 - max_iter
- Method:
 - fit
 - predict

Reference: <https://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html> (<https://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>)

Using a pre-built library yields no credit. You have to write everything from scratch

```
In [18]: import math
import numpy as np

def euclidean(point, data):
    return np.sqrt(np.sum((point - data)**2, axis=1))

class My_KMeans:

    l= []

    def __init__(self, n_clusters = 3, max_iter = 100):

        self.n_clusters = n_clusters
        self.max_iter = max_iter
        self.centroids =0

    def fit(self, X_train):

        min_, max_ = np.min(X_train, axis=0), np.max(X_train, axis=0)
        self.centroids = [np.random.uniform(min_, max_) for _ in range(self.n_clusters)]
        iteration = 0

        prev_centroids = None

        while np.not_equal(self.centroids, prev_centroids).any() and iteration < self.max_iter:

            sorted_points = [[] for _ in range(self.n_clusters)]
            for x in X_train:
                dists = euclidean(x, self.centroids)
                centroid_idx = np.argmin(dists)
                sorted_points[centroid_idx].append(x)
            prev_centroids = self.centroids
            self.centroids = [np.mean(cluster, axis=0) for cluster in sorted_points]
            for i, centroid in enumerate(self.centroids):
                if np.isnan(centroid).any():
                    self.centroids[i] = prev_centroids[i]
            iteration += 1

    def predict(self, X):
        centroids = []
        centroid_idx = []
        for x in X:
```

```
dists = euclidean(x, self.centroids)
centroid_idx = np.argmin(dists)
t = self.centroids[centroid_idx]
centroids.append(t)
centroid_idxs.append(centroid_idx)
cluster_centers_ = centroid_idxs
self.l = centroid_idxs
return centroid_idxs
```

Run the code

```
In [19]: learn.preprocessing import StandardScaler
matplotlib.pyplot as plt
seaborn sns

standardize the data
StandardScaler().fit_transform(df)


local implementation of kmeans
k = My_KMeans(3,100)
k.fit(X_std)
y = k.predict(X_std)

, 1]: X_std dataset with all rows and the second column (time_study in this case)
, 2]: X_std dataset with all rows and the third column (marks in this case)

cluster elements
scatter(x=X_std[:, 1], y=X_std[:, 2], c=y_kmeans, s=50, cmap='viridis')

cluster centers
scatter([x for x, _, _ in kmeans.centroids], [y for _, _, y in kmeans.centroids], marker='*', color='r', s=200)
```

Out[19]: <matplotlib.collections.PathCollection at 0x7f792a020490>



Your analysis about the plot

According to my inspection of the previous plot, both the points and clusters, as well as the centroid that was created, are not precise. Even if the model appears to be functioning OK, a few corrections should still need to be made. However, the model that is now being created may be applied since the outcomes are adequate.

Type *Markdown* and LaTeX: α^2

(Bonus 20pts) Modify your algorithm to k-means++ with a smart initialization, and repeat the task above

```

In [20]: def get_kmeans_pp_centroids(X1,k = 5):
    centroids = []
    for i in range(k):
        centroids.append([X1[10,0],X1[10,1],X1[10,2]])
    i = 1
    d = []
    while i <= k:
        r = centroids[i-2 if i-2 >=0 else 0]
        maxd = 0
        for j in X1:
            d = np.sqrt((j[1] - r[1])**2 +(j[2] - r[2])**2)
            if d > maxd:
                centroids[i-1] = [j[0],j[1],j[2]]
                maxd = d
        i+=1
    min_, max_ = np.min(X1, axis=0), np.max(X1, axis=0)
    iteration = 0
    prev_centroids = None

    while np.not_equal(centroids, prev_centroids).any() and iteration < k:

        sorted_points = [[] for _ in range(k)]

        for x in X1:
            dists = euclidean(x, centroids)
            centroid_idx = np.argmin(dists)
            sorted_points[centroid_idx].append(x)

        prev_centroids = centroids
        centroids = [np.mean(cluster, axis=0) for cluster in sorted_points]
        for i, centroid in enumerate(centroids):
            if np.isnan(centroid).any():
                centroids[i] = prev_centroids[i]
        iteration += 1
    return centroids

def predict(X,c):
    centroids = []
    centroid_idx = []
    for x in X:
        dists = euclidean(x,c)
        centroid_idx = np.argmin(dists)

```

```

    t = c[centroid_idx]
    centroids.append(t)
    centroid_idxs.append(centroid_idx)
    return centroid_idxs

```

```

In [21]: centroids = get_kmeans_pp_centroids(X_std, k = 3)
y_kmeans = predict(X_std, centroids)

# Plot cluster elements
plt.scatter(x=X_std[:, 1], y=X_std[:, 2], c=y_kmeans, s=50, cmap='viridis')

# Plot cluster centers
plt.scatter([x for x, _, _ in centroids], [y for _, _, y in centroids], marker='*', color='r', s=200, alp

```


/Users/roja/opt/anaconda3/lib/python3.9/site-packages/numpy/core/fromnumeric.py:3419: RuntimeWarning: Mean of empty slice.

```

    return _methods._mean(a, axis=axis, dtype=dtype,
/Users/roja/opt/anaconda3/lib/python3.9/site-packages/numpy/core/_methods.py:188: RuntimeWarning: invalid value encountered in double_scalars
    ret = ret.dtype.type(ret / rcount)

```

Out[21]: <matplotlib.collections.PathCollection at 0x7f792a1ea280>



In []:

